


The training loss does not predict the benchmark

A. Training loss falls monotonically



B. Effective resolution gets WORSE — in both runs

